

ModelResults

March 7, 2025

```
[3]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

When we look at the data, we realize that the target columns are actually stored in a separate file. To fix this, we can just merge the two since they share ids into one dataframe. We will combine functional with functional needs to repairs so this will be a binary problem as opposed to a ternary one.

```
[4]: df_labels = pd.read_csv('../data/training_set_labels.csv', header=0)
df_labels.replace({'functional needs repair': 'functional'}, inplace=True)
df_labels['status_group'].value_counts(normalize=True)
```

```
[4]: functional      0.615758
non functional     0.384242
Name: status_group, dtype: float64
```

```
[5]: df_values = pd.read_csv('../data/training_set_values.csv', header=0)
df_values
```

```
[5]:
```

	id	amount_tsh	date_recorded	funder	gps_height	\
0	69572	6000.0	2011-03-14	Roman	1390	
1	8776	0.0	2013-03-06	Grumeti	1399	
2	34310	25.0	2013-02-25	Lottery Club	686	
3	67743	0.0	2013-01-28	Unicef	263	
4	19728	0.0	2011-07-13	Action In A	0	
...	
59395	60739	10.0	2013-05-03	Germany Republi	1210	
59396	27263	4700.0	2011-05-07	Cefa-njombe	1212	
59397	37057	0.0	2011-04-11	NaN	0	
59398	31282	0.0	2011-03-08	Malec	0	
59399	26348	0.0	2011-03-23	World Bank	191	

	installer	longitude	latitude	wpt_name	num_private	\
0	Roman	34.938093	-9.856322	none	0	
1	GRUMETI	34.698766	-2.147466	Zahanati	0	

2	World vision	37.460664	-3.821329	Kwa Mahundi	0
3	UNICEF	38.486161	-11.155298	Zahanati Ya Nanyumbu	0
4	Artisan	31.130847	-1.825359	Shuleni	0
...
59395	CES	37.169807	-3.253847	Area Three Namba 27	0
59396	Cefa	35.249991	-9.070629	Kwa Yahona Kuvala	0
59397	NaN	34.017087	-8.750434	Mashine	0
59398	Musa	35.861315	-6.378573	Mshoro	0
59399	World	38.104048	-6.747464	Kwa Mzee Lugawa	0

	...	payment_type	water_quality	quality_group	quantity	\
0	...	annually	soft	good	enough	
1	...	never pay	soft	good	insufficient	
2	...	per bucket	soft	good	enough	
3	...	never pay	soft	good	dry	
4	...	never pay	soft	good	seasonal	

...
59395	...	per bucket	soft	good	enough	
59396	...	annually	soft	good	enough	
59397	...	monthly	fluoride	fluoride	enough	
59398	...	never pay	soft	good	insufficient	
59399	...	on failure	salty	salty	enough	

	quantity_group	source	source_type	\
0	enough	spring	spring	
1	insufficient	rainwater harvesting	rainwater harvesting	
2	enough	dam	dam	
3	dry	machine dbh	borehole	
4	seasonal	rainwater harvesting	rainwater harvesting	
...
59395	enough	spring	spring	
59396	enough	river	river/lake	
59397	enough	machine dbh	borehole	
59398	insufficient	shallow well	shallow well	
59399	enough	shallow well	shallow well	

	source_class	waterpoint_type	waterpoint_type_group
0	groundwater	communal standpipe	communal standpipe
1	surface	communal standpipe	communal standpipe
2	surface	communal standpipe multiple	communal standpipe
3	groundwater	communal standpipe multiple	communal standpipe
4	surface	communal standpipe	communal standpipe
...
59395	groundwater	communal standpipe	communal standpipe
59396	surface	communal standpipe	communal standpipe
59397	groundwater	hand pump	hand pump
59398	groundwater	hand pump	hand pump

59399 groundwater hand pump hand pump

[59400 rows x 40 columns]

Since the data is stored in two separate csv, we will combine them into one. Then we'll combine the target columns with the dataframes.

```
[6]: df_val_test = pd.read_csv('../data/test_set.csv', header=0)
df_val_test
df_values = pd.concat([df_values, df_val_test], ignore_index=True)
df_values = df_values.merge(df_labels, left_on='id', right_on='id')
```

```
[7]: df_values
```

```
[7]:
```

	id	amount_tsh	date_recorded	funder	gps_height	\
0	69572	6000.0	2011-03-14	Roman	1390	
1	8776	0.0	2013-03-06	Grumeti	1399	
2	34310	25.0	2013-02-25	Lottery Club	686	
3	67743	0.0	2013-01-28	Unicef	263	
4	19728	0.0	2011-07-13	Action In A	0	
...	
59395	60739	10.0	2013-05-03	Germany Republi	1210	
59396	27263	4700.0	2011-05-07	Cefa-njombe	1212	
59397	37057	0.0	2011-04-11	NaN	0	
59398	31282	0.0	2011-03-08	Malec	0	
59399	26348	0.0	2011-03-23	World Bank	191	

	installer	longitude	latitude	wpt_name	num_private	\
0	Roman	34.938093	-9.856322	none	0	
1	GRUMETI	34.698766	-2.147466	Zahanati	0	
2	World vision	37.460664	-3.821329	Kwa Mahundi	0	
3	UNICEF	38.486161	-11.155298	Zahanati Ya Nanyumbu	0	
4	Artisan	31.130847	-1.825359	Shuleni	0	
...	
59395	CES	37.169807	-3.253847	Area Three Namba 27	0	
59396	Cefa	35.249991	-9.070629	Kwa Yahona Kuvala	0	
59397	NaN	34.017087	-8.750434	Mashine	0	
59398	Musa	35.861315	-6.378573	Mshoro	0	
59399	World	38.104048	-6.747464	Kwa Mzee Lugawa	0	

	...	water_quality	quality_group	quantity	quantity_group	\
0	...	soft	good	enough	enough	
1	...	soft	good	insufficient	insufficient	
2	...	soft	good	enough	enough	
3	...	soft	good	dry	dry	
4	...	soft	good	seasonal	seasonal	
...	
59395	...	soft	good	enough	enough	

59396	...	soft	good	enough	enough
59397	...	fluoride	fluoride	enough	enough
59398	...	soft	good	insufficient	insufficient
59399	...	salty	salty	enough	enough

		source	source_type	source_class	\
0		spring	spring	groundwater	
1	rainwater	harvesting	rainwater harvesting	surface	
2		dam	dam	surface	
3		machine dbh	borehole	groundwater	
4	rainwater	harvesting	rainwater harvesting	surface	
...		
59395		spring	spring	groundwater	
59396		river	river/lake	surface	
59397		machine dbh	borehole	groundwater	
59398		shallow well	shallow well	groundwater	
59399		shallow well	shallow well	groundwater	

		waterpoint_type	waterpoint_type_group	status_group
0		communal standpipe	communal standpipe	functional
1		communal standpipe	communal standpipe	functional
2	communal	standpipe multiple	communal standpipe	functional
3	communal	standpipe multiple	communal standpipe	non functional
4		communal standpipe	communal standpipe	functional
...	
59395		communal standpipe	communal standpipe	functional
59396		communal standpipe	communal standpipe	functional
59397		hand pump	hand pump	functional
59398		hand pump	hand pump	functional
59399		hand pump	hand pump	functional

[59400 rows x 41 columns]

There are a lot of columns with non numerical entries. This means we might have to one hot encode them, however with how many columns and distinct entries there are, it might be too many factors. We will likely need to drop some of these columns. First lets plot of graphs with only the numeric columns

We also should take a sample of the data for this graphing since there is a lot of data and it would take a long time to graph each time.

```
[8]: df_numeric = df_values.select_dtypes(include=np.number).merge(df_labels,
    ↪left_on='id', right_on='id')
    #Take a sample of a quarter of the data, random state for reproducibility
    df_num_sample = df_numeric.sample(frac=0.10, random_state= 5)
    df_num_sample.drop(columns=['id'], inplace=True)
```

```
[9]: df_small = df_num_sample[['amount_tsh', 'gps_height', 'population', '
    ↪ 'construction_year', 'region_code', 'status_group']]
sns.pairplot(hue='status_group', data=df_small)
#sns.pairplot(hue = 'status_group', data= df_num_sample)
```


```
[9]: <seaborn.axisgrid.PairGrid at 0x29cb4816460>
```



Here we'll look more closely at some specific examples

```
[10]: sns.scatterplot(data= df_num_sample, x = 'region_code', y='gps_height', hue = '
    ↪ 'status_group')
```

```
[10]: <AxesSubplot:xlabel='region_code', ylabel='gps_height'>
```



```
[11]: df_values.columns
```

```
[11]: Index(['id', 'amount_tsh', 'date_recorded', 'funder', 'gps_height',
            'installer', 'longitude', 'latitude', 'wpt_name', 'num_private',
            'basin', 'subvillage', 'region', 'region_code', 'district_code', 'lga',
            'ward', 'population', 'public_meeting', 'recorded_by',
            'scheme_management', 'scheme_name', 'permit', 'construction_year',
            'extraction_type', 'extraction_type_group', 'extraction_type_class',
            'management', 'management_group', 'payment', 'payment_type',
            'water_quality', 'quality_group', 'quantity', 'quantity_group',
            'source', 'source_type', 'source_class', 'waterpoint_type',
            'waterpoint_type_group', 'status_group'],
            dtype='object')
```

```
[12]: good_qual_df = df_values[(df_values['quality_group']=='good') &
                                ↪(df_values['amount_tsh']<100)]
good_qual_df
```

```
[12]:
```

	id	amount_tsh	date_recorded	funder	gps_height	\
1	8776	0.0	2013-03-06	Grumeti	1399	
2	34310	25.0	2013-02-25	Lottery Club	686	
3	67743	0.0	2013-01-28	Unicef	263	
4	19728	0.0	2011-07-13	Action In A	0	
6	19816	0.0	2012-10-01	Dwsp	0	

...
59391	44885	0.0	2013-08-03	Government Of Tanzania	540	
59392	40607	0.0	2011-04-15	Government Of Tanzania	0	
59393	48348	0.0	2012-10-27	Private	0	
59395	60739	10.0	2013-05-03	Germany Republi	1210	
59398	31282	0.0	2011-03-08	Malec	0	

	installer	longitude	latitude	wpt_name	num_private	\
1	GRUMETI	34.698766	-2.147466	Zahanati	0	
2	World vision	37.460664	-3.821329	Kwa Mahundi	0	
3	UNICEF	38.486161	-11.155298	Zahanati Ya Nanyumbu	0	
4	Artisan	31.130847	-1.825359	Shuleni	0	
6	DWSP	33.362410	-3.766365	Kwa Ngomho	0	

...
59391	Government	38.044070	-4.272218	Kwa	0	
59392	Government	33.009440	-8.520888	Benard Charles	0	
59393	Private	33.866852	-4.287410	Kwa Peter	0	
59395	CES	37.169807	-3.253847	Area Three Namba 27	0	
59398	Musa	35.861315	-6.378573	Mshoro	0	

	...	water_quality	quality_group	quantity	quantity_group	\
1	...	soft	good	insufficient	insufficient	
2	...	soft	good	enough	enough	
3	...	soft	good	dry	dry	
4	...	soft	good	seasonal	seasonal	
6	...	soft	good	enough	enough	

...
59391	...	soft	good	enough	enough	
59392	...	soft	good	enough	enough	
59393	...	soft	good	insufficient	insufficient	
59395	...	soft	good	enough	enough	
59398	...	soft	good	insufficient	insufficient	

		source	source_type	source_class	\
1	rainwater harvesting	rainwater harvesting	surface		
2	dam	dam	surface		
3	machine dbh	borehole	groundwater		
4	rainwater harvesting	rainwater harvesting	surface		
6	machine dbh	borehole	groundwater		

...
59391	river	river/lake	surface			
59392	spring	spring	groundwater			
59393	dam	dam	surface			
59395	spring	spring	groundwater			
59398	shallow well	shallow well	groundwater			


	waterpoint_type	waterpoint_type_group	status_group
--	-----------------	-----------------------	--------------

1	communal	standpipe	communal	standpipe	functional
2	communal	standpipe multiple	communal	standpipe	functional
3	communal	standpipe multiple	communal	standpipe	non functional
4	communal	standpipe	communal	standpipe	functional
6		hand pump		hand pump	non functional
...	
59391	communal	standpipe	communal	standpipe	non functional
59392	communal	standpipe	communal	standpipe	non functional
59393		other		other	functional
59395	communal	standpipe	communal	standpipe	functional
59398		hand pump		hand pump	functional

[40493 rows x 41 columns]

```
[13]: sns.scatterplot(data=good_qual_df, x = 'amount_tsh', y='quality_group', hue = 'status_group')
```

```
[13]: <AxesSubplot:xlabel='amount_tsh', ylabel='quality_group'>
```



We can see that in most of the cases, the functional and non-functional had no clear separations. So the numerical columns don't seem to be good at determining whether a well is functional or not.

1 Baseline Model

With just the numerical categories, there does not seem to be any columns that do a particularly good job at correlating to a well being either function or non functional.

What we'll do for now is just to create a baseline model without any preprocessing to see what the accuracy score will be.

```
[14]: #Drop the id column now that we don't need it
df_values.drop(columns=['id'], inplace=True)
df_values['status_group'].replace({'functional': 1, 'non functional': 0},
    ↪inplace=True)

[15]: from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split, cross_validate
from sklearn.metrics import accuracy_score, recall_score

X = df_numeric.drop(columns=['status_group'])
y = df_values['status_group']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25,
    ↪random_state=42)

model = LogisticRegression(max_iter=10000, solver='liblinear')
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
print(accuracy_score(y_test, y_pred))
print(f"Specificity(Recall) Score: {recall_score(y_test, y_pred, pos_label=0)}")
print(f"Recall Score: {recall_score(y_test, y_pred)}")
```

0.6174410774410775

Specificity(Recall) Score: 0.10936949630151462

Recall Score: 0.9319668556476232

Here we are creating lists to store the accuracy, specificity and recall score so we can graph a line chart at the end that measures how they change

```
[16]: list_acc = []
list_spec = []
list_rec = []
list_acc.append(accuracy_score(y_test, y_pred))
list_spec.append(recall_score(y_test, y_pred, pos_label=0))
list_rec.append(recall_score(y_test, y_pred))
```


```
[17]: df_values['status_group'].value_counts()
```

```
[17]: 1    36576
0     22824
Name: status_group, dtype: int64
```

Without any preprocessing done, the model does a pretty bad job at predicting the results, with a large amount of false positives and a not so small false negative as well.

```
[18]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, \
      ↪ recall_score, precision_score
fig, ax = plt.subplots(figsize=(10,8))
fig.set_facecolor('grey')
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(cm)
disp.plot(ax=ax)
```

```
[18]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x29cb7b4abe0>
```



Next, we'll try to look at the coefficients to see which columns have a high correlation.

```
[19]: coefficients = model.coef_[0]
      odds_ratios = np.exp(coefficients)
```

```
feature_importance = pd.DataFrame({'Feature': X.columns, 'Coefficients':  
    ↪coefficients, 'Odds Ratios': odds_ratios})  
feature_importance
```

```
[19]:
```

	Feature	Coefficients	Odds Ratios
0	id	0.000002	1.000002
1	amount_tsh	0.000267	1.000267
2	gps_height	0.000386	1.000386
3	longitude	0.010821	1.010879
4	latitude	-0.000632	0.999368
5	num_private	0.000817	1.000818
6	region_code	-0.006937	0.993087
7	district_code	-0.000543	0.999457
8	population	0.000102	1.000102
9	construction_year	-0.000157	0.999843

Here we can see that most of the numerical data has a negative coefficient, aside from longitude, num private, region code, and construction year, most of which have a near 0 coefficient. We'll also look at the categorical columns and see if we can drop any that are not particularly useful.

```
[20]: df_values['extraction_type_group'].value_counts()
```

```
[20]: gravity                26780  
nira/tanira                 8154  
other                       6430  
submersible                 6179  
swn 80                      3670  
mono                       2865  
india mark ii              2400  
afridev                    1770  
rope pump                   451  
other handpump              364  
other motorpump            122  
wind-powered               117  
india mark iii              98  
Name: extraction_type_group, dtype: int64
```

```
[21]: #A fair amount of these columns are essentially duplicates, have the same data,  
    ↪as other columns or similar  
#PUT BACK water_quality, management, source, district_code, extraction_type  
non_important_columns = ['wpt_name', 'region_code', 'subvillage',  
    ↪'extraction_type', 'payment', 'quantity', 'waterpoint_type', 'source_type',  
    ↪'scheme_name', 'lga', 'ward']  
prep_df_values = df_values.drop(columns=non_important_columns)  
prep_df_values
```

[21]:

	amount_tsh	date_recorded	funder	gps_height	installer \
0	6000.0	2011-03-14	Roman	1390	Roman
1	0.0	2013-03-06	Grumeti	1399	GRUMETI
2	25.0	2013-02-25	Lottery Club	686	World vision
3	0.0	2013-01-28	Unicef	263	UNICEF
4	0.0	2011-07-13	Action In A	0	Artisan
...
59395	10.0	2013-05-03	Germany Republi	1210	CES
59396	4700.0	2011-05-07	Cefa-njombe	1212	Cefa
59397	0.0	2011-04-11	NaN	0	NaN
59398	0.0	2011-03-08	Malec	0	Musa
59399	0.0	2011-03-23	World Bank	191	World

	longitude	latitude	num_private	basin \
0	34.938093	-9.856322	0	Lake Nyasa
1	34.698766	-2.147466	0	Lake Victoria
2	37.460664	-3.821329	0	Pangani
3	38.486161	-11.155298	0	Ruvuma / Southern Coast
4	31.130847	-1.825359	0	Lake Victoria
...
59395	37.169807	-3.253847	0	Pangani
59396	35.249991	-9.070629	0	Rufiji
59397	34.017087	-8.750434	0	Rufiji
59398	35.861315	-6.378573	0	Rufiji
59399	38.104048	-6.747464	0	Wami / Ruvu

	region	...	management	management_group	payment_type \
0	Iringa	...	vwc	user-group	annually
1	Mara	...	wug	user-group	never pay
2	Manyara	...	vwc	user-group	per bucket
3	Mtwara	...	vwc	user-group	never pay
4	Kagera	...	other	other	never pay
...
59395	Kilimanjaro	...	water board	user-group	per bucket
59396	Iringa	...	vwc	user-group	annually
59397	Mbeya	...	vwc	user-group	monthly
59398	Dodoma	...	vwc	user-group	never pay
59399	Morogoro	...	vwc	user-group	on failure

	water_quality	quality_group	quantity_group	source \
0	soft	good	enough	spring
1	soft	good	insufficient	rainwater harvesting
2	soft	good	enough	dam
3	soft	good	dry	machine dbh
4	soft	good	seasonal	rainwater harvesting
...
59395	soft	good	enough	spring

59396	soft	good	enough	river
59397	fluoride	fluoride	enough	machine dbh
59398	soft	good	insufficient	shallow well
59399	salty	salty	enough	shallow well

	source_class	waterpoint_type_group	status_group
0	groundwater	communal standpipe	1
1	surface	communal standpipe	1
2	surface	communal standpipe	1
3	groundwater	communal standpipe	0
4	surface	communal standpipe	1
...
59395	groundwater	communal standpipe	1
59396	surface	communal standpipe	1
59397	groundwater	hand pump	1
59398	groundwater	hand pump	1
59399	groundwater	hand pump	1

[59400 rows x 29 columns]

We can also combine date_recorded with construction year by subtracting the two and storing them as a year

```
[22]: from sklearn.preprocessing import OneHotEncoder
from datetime import datetime
prep_df_values['date_recorded'] = prep_df_values['date_recorded'].apply(lambda row: row[:4])
prep_df_values['years_active'] = prep_df_values['date_recorded'].astype(int) - prep_df_values['construction_year'].astype(int)

#Convert years with 0 as construction year to simply 0
prep_df_values['years_active'] = prep_df_values['years_active'].apply(lambda year: year if year < 100 else np.nan)
#TEMP TRYING
prep_df_values[['amount_tsh', 'gps_height', 'longitude', 'latitude', 'population']] = prep_df_values[['amount_tsh', 'gps_height', 'longitude', 'latitude', 'population']].replace(0, np.nan)
prep_df_values.replace('unknown', np.nan)
#Now drop the two other columns
prep_df_values.drop(columns=['date_recorded', 'construction_year'], inplace=True)

prep_df_values
```

```
[22]: amount_tsh      funder  gps_height  installer  longitude \
0      6000.0      Roman    1390.0      Roman    34.938093
1       NaN    Grumeti    1399.0    GRUMETI    34.698766
```

2	25.0	Lottery Club	686.0	World vision	37.460664
3	NaN	Unicef	263.0	UNICEF	38.486161
4	NaN	Action In A	NaN	Artisan	31.130847
...
59395	10.0	Germany Republi	1210.0	CES	37.169807
59396	4700.0	Cefa-njombe	1212.0	Cefa	35.249991
59397	NaN	NaN	NaN	NaN	34.017087
59398	NaN	Malec	NaN	Musa	35.861315
59399	NaN	World Bank	191.0	World	38.104048

	latitude	num_private	basin	region \
0	-9.856322	0	Lake Nyasa	Iringa
1	-2.147466	0	Lake Victoria	Mara
2	-3.821329	0	Pangani	Manyara
3	-11.155298	0	Ruvuma / Southern Coast	Mtwara
4	-1.825359	0	Lake Victoria	Kagera

...
59395	-3.253847	0	Pangani	Kilimanjaro
59396	-9.070629	0	Rufiji	Iringa
59397	-8.750434	0	Rufiji	Mbeya
59398	-6.378573	0	Rufiji	Dodoma
59399	-6.747464	0	Wami / Ruvu	Morogoro

	district_code	...	management_group	payment_type	water_quality \
0	5	...	user-group	annually	soft
1	2	...	user-group	never pay	soft
2	4	...	user-group	per bucket	soft
3	63	...	user-group	never pay	soft
4	1	...	other	never pay	soft
...
59395	5	...	user-group	per bucket	soft
59396	4	...	user-group	annually	soft
59397	7	...	user-group	monthly	fluoride
59398	4	...	user-group	never pay	soft
59399	2	...	user-group	on failure	salty

	quality_group	quantity_group	source	source_class \
0	good	enough	spring	groundwater
1	good	insufficient	rainwater harvesting	surface
2	good	enough	dam	surface
3	good	dry	machine dbh	groundwater
4	good	seasonal	rainwater harvesting	surface
...
59395	good	enough	spring	groundwater
59396	good	enough	river	surface
59397	fluoride	enough	machine dbh	groundwater
59398	good	insufficient	shallow well	groundwater

59399	salty	enough	shallow well	groundwater
-------	-------	--------	--------------	-------------

	waterpoint_type_group	status_group	years_active
0	communal standpipe	1	12.0
1	communal standpipe	1	3.0
2	communal standpipe	1	4.0
3	communal standpipe	0	27.0
4	communal standpipe	1	NaN
...
59395	communal standpipe	1	14.0
59396	communal standpipe	1	15.0
59397	hand pump	1	NaN
59398	hand pump	1	NaN
59399	hand pump	1	9.0

[59400 rows x 28 columns]

Here we will create a dataframe for specifically the categorical columns

```
[23]: df_categoricals = prep_df_values.select_dtypes(exclude=np.number)
df_categoricals
```

```
[23]:
```

	funder	installer	basin	region \
0	Roman	Roman	Lake Nyasa	Iringa
1	Grumeti	GRUMETI	Lake Victoria	Mara
2	Lottery Club	World vision	Pangani	Manyara
3	Unicef	UNICEF	Ruvuma / Southern Coast	Mtwara
4	Action In A	Artisan	Lake Victoria	Kagera
...
59395	Germany Republi	CES	Pangani	Kilimanjaro
59396	Cefa-njombe	Cefa	Rufiji	Iringa
59397	NaN	NaN	Rufiji	Mbeya
59398	Malec	Musa	Rufiji	Dodoma
59399	World Bank	World	Wami / Ruvu	Morogoro

	public_meeting	recorded_by	scheme_management	permit \
0	True	GeoData Consultants Ltd	VWC	False
1	NaN	GeoData Consultants Ltd	Other	True
2	True	GeoData Consultants Ltd	VWC	True
3	True	GeoData Consultants Ltd	VWC	True
4	True	GeoData Consultants Ltd	NaN	True
...
59395	True	GeoData Consultants Ltd	Water Board	True
59396	True	GeoData Consultants Ltd	VWC	True
59397	True	GeoData Consultants Ltd	VWC	False
59398	True	GeoData Consultants Ltd	VWC	True
59399	True	GeoData Consultants Ltd	VWC	True

	extraction_type_group	extraction_type_class	management	\
0	gravity	gravity	vwc	
1	gravity	gravity	wug	
2	gravity	gravity	vwc	
3	submersible	submersible	vwc	
4	gravity	gravity	other	
...	
59395	gravity	gravity	water board	
59396	gravity	gravity	vwc	
59397	swn 80	handpump	vwc	
59398	nira/tanira	handpump	vwc	
59399	nira/tanira	handpump	vwc	

	management_group	payment_type	water_quality	quality_group	\
0	user-group	annually	soft	good	
1	user-group	never pay	soft	good	
2	user-group	per bucket	soft	good	
3	user-group	never pay	soft	good	
4	other	never pay	soft	good	
...	
59395	user-group	per bucket	soft	good	
59396	user-group	annually	soft	good	
59397	user-group	monthly	fluoride	fluoride	
59398	user-group	never pay	soft	good	
59399	user-group	on failure	salty	salty	

	quantity_group	source	source_class	waterpoint_type_group
0	enough	spring	groundwater	communal standpipe
1	insufficient	rainwater harvesting	surface	communal standpipe
2	enough	dam	surface	communal standpipe
3	dry	machine dbh	groundwater	communal standpipe
4	seasonal	rainwater harvesting	surface	communal standpipe
...
59395	enough	spring	groundwater	communal standpipe
59396	enough	river	surface	communal standpipe
59397	enough	machine dbh	groundwater	hand pump
59398	insufficient	shallow well	groundwater	hand pump
59399	enough	shallow well	groundwater	hand pump

[59400 rows x 19 columns]

2 Building up Logistic Regression Model

Now is the time to one hot encode the categorical columns to see if there is a correlation. For any columns that end up being encoded into hundreds or even thousands of columns, they aren't worth considering.


```
[24]: y = prep_df_values['status_group']
X = prep_df_values[['amount_tsh', 'gps_height', 'longitude', 'latitude',
↳ 'num_private', 'population', 'years_active', 'district_code',
↳ 'basin', 'region', 'management_group', 'quality_group',
↳ 'permit', 'water_quality', 'management', 'source', 'extraction_type_class',
↳ 'extraction_type_group',
↳ 'waterpoint_type_group', 'quantity_group', 'payment_type']]
X
```

```
[24]:
```

	amount_tsh	gps_height	longitude	latitude	num_private	population	\
0	6000.0	1390.0	34.938093	-9.856322	0	109.0	
1	NaN	1399.0	34.698766	-2.147466	0	280.0	
2	25.0	686.0	37.460664	-3.821329	0	250.0	
3	NaN	263.0	38.486161	-11.155298	0	58.0	
4	NaN	NaN	31.130847	-1.825359	0	NaN	
...	
59395	10.0	1210.0	37.169807	-3.253847	0	125.0	
59396	4700.0	1212.0	35.249991	-9.070629	0	56.0	
59397	NaN	NaN	34.017087	-8.750434	0	NaN	
59398	NaN	NaN	35.861315	-6.378573	0	NaN	
59399	NaN	191.0	38.104048	-6.747464	0	150.0	

	years_active	district_code	basin	region	...	\
0	12.0	5	Lake Nyasa	Iringa	...	
1	3.0	2	Lake Victoria	Mara	...	
2	4.0	4	Pangani	Manyara	...	
3	27.0	63	Ruvuma / Southern Coast	Mtwara	...	
4	NaN	1	Lake Victoria	Kagera	...	
...	
59395	14.0	5	Pangani	Kilimanjaro	...	
59396	15.0	4	Rufiji	Iringa	...	
59397	NaN	7	Rufiji	Mbeya	...	
59398	NaN	4	Rufiji	Dodoma	...	
59399	9.0	2	Wami / Ruvu	Morogoro	...	

	quality_group	permit	water_quality	management	source	\
0	good	False	soft	vwc	spring	
1	good	True	soft	wug	rainwater harvesting	
2	good	True	soft	vwc	dam	
3	good	True	soft	vwc	machine dbh	
4	good	True	soft	other	rainwater harvesting	
...	
59395	good	True	soft	water board	spring	
59396	good	True	soft	vwc	river	
59397	fluoride	False	fluoride	vwc	machine dbh	
59398	good	True	soft	vwc	shallow well	
59399	salty	True	salty	vwc	shallow well	

	extraction_type_class	extraction_type_group	waterpoint_type_group	\
0	gravity	gravity	communal	standpipe
1	gravity	gravity	communal	standpipe
2	gravity	gravity	communal	standpipe
3	submersible	submersible	communal	standpipe
4	gravity	gravity	communal	standpipe
...
59395	gravity	gravity	communal	standpipe
59396	gravity	gravity	communal	standpipe
59397	handpump	swn 80		hand pump
59398	handpump	nira/tanira		hand pump
59399	handpump	nira/tanira		hand pump

	quantity_group	payment_type
0	enough	annually
1	insufficient	never pay
2	enough	per bucket
3	dry	never pay
4	seasonal	never pay
...
59395	enough	per bucket
59396	enough	annually
59397	enough	monthly
59398	insufficient	never pay
59399	enough	on failure

[59400 rows x 21 columns]

Here we see that with the new relevant numerical columns and the one hot encoded columns that are worth exploring, we actually increased our accuracy metric to 78 from around 65.

```
[25]: from sklearn.preprocessing import StandardScaler
from sklearn.experimental import enable_iterative_imputer
from sklearn.impute import SimpleImputer, IterativeImputer
from imblearn.over_sampling import SMOTE
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25,
↳ random_state=41)

num_imp = SimpleImputer(strategy='median')
cat_imp = SimpleImputer(strategy='most_frequent')
#Impute numeric and categorical columns for train and test set
X_train_num = X_train[['amount_tsh', 'gps_height', 'longitude', 'latitude',
↳ 'num_private', 'population', 'years_active', 'district_code']]
```

```

X_train_cat = X_train[['basin', 'region', 'management_group', 'quality_group',
    ↳ 'permit', 'waterpoint_type_group', 'quantity_group', 'payment_type',
    ↳ 'water_quality', 'management', 'source', 'extraction_type_group',
    ↳ 'extraction_type_class']]
X_test_num = X_test[['amount_tsh', 'gps_height', 'longitude', 'latitude',
    ↳ 'num_private', 'population', 'years_active', 'district_code']]
X_test_cat = X_test[['basin', 'region', 'management_group', 'quality_group',
    ↳ 'permit', 'waterpoint_type_group', 'quantity_group', 'payment_type',
    ↳ 'water_quality', 'management', 'source', 'extraction_type_group',
    ↳ 'extraction_type_class']]

#Fit transform for train, transform test
num_arr = num_imp.fit_transform(X_train_num)
cat_arr = cat_imp.fit_transform(X_train_cat)
num_t_arr = num_imp.transform(X_test_num)
cat_t_arr = cat_imp.transform(X_test_cat)

#One Hot Encode Train
X_train_imp = pd.concat([pd.DataFrame(num_arr, columns=X_train_num.columns,
    ↳ index= X_train_num.index), pd.DataFrame(cat_arr, columns=X_train_cat.
    ↳ columns, index=X_train_cat.index)], axis=1)
X_train_imp = pd.get_dummies(X_train_imp,
    columns=['basin', 'region', 'management_group',
    ↳ 'quality_group', 'permit', 'waterpoint_type_group', 'quantity_group',
    ↳ 'payment_type', 'water_quality', 'management',
    ↳ 'source', 'extraction_type_group', 'extraction_type_class'], drop_first=True,
    ↳ dtype=int)

#Test
X_test_imp = pd.concat([pd.DataFrame(num_t_arr, columns=X_test_num.columns,
    ↳ index= X_test_num.index), pd.DataFrame(cat_t_arr, columns=X_test_cat.
    ↳ columns, index=X_test_cat.index)], axis=1)
X_test_imp = pd.get_dummies(X_test_imp,
    columns=['basin', 'region', 'management_group',
    ↳ 'quality_group', 'permit', 'waterpoint_type_group', 'quantity_group',
    ↳ 'payment_type', 'water_quality', 'management',
    ↳ 'source', 'extraction_type_group', 'extraction_type_class'], drop_first=True,
    ↳ dtype=int)

#Oversample
smote = SMOTE(random_state=42, sampling_strategy= 1)
X_train_samp, y_train_samp = smote.fit_resample(X_train_imp, y_train)
#X_train_samp, y_train_samp = X_train_imp, y_train

#Fit model

```

```

model = LogisticRegression(solver='liblinear', max_iter=10000)
model.fit(X_train_samp, y_train_samp)

y_train_pred = model.predict(X_train_samp)
y_pred = model.predict(X_test_imp)

print(f"Train Acc: {accuracy_score(y_train_samp, y_train_pred)}\nTest Acc: {accuracy_score(y_test, y_pred)}")
print(f"Specificity(Recall) Score: {recall_score(y_test, y_pred, pos_label=0)}")
print(f"Recall Score: {recall_score(y_test, y_pred)}")

```

Train Acc: 0.8042949020747793

Test Acc: 0.7893602693602694

Specificity(Recall) Score: 0.6374460742018981

Recall Score: 0.8865819988956378


We can see that we have increased the number of true negatives, but we still have a large amount of the false metrics. In fact, we got a little more false negatives.

```

[26]: cm = confusion_matrix(y_test, y_pred)
      disp = ConfusionMatrixDisplay(cm)
      disp.plot()

```

[26]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x29cb649a7c0>



3 Hyperparameter Tuning (RAW)

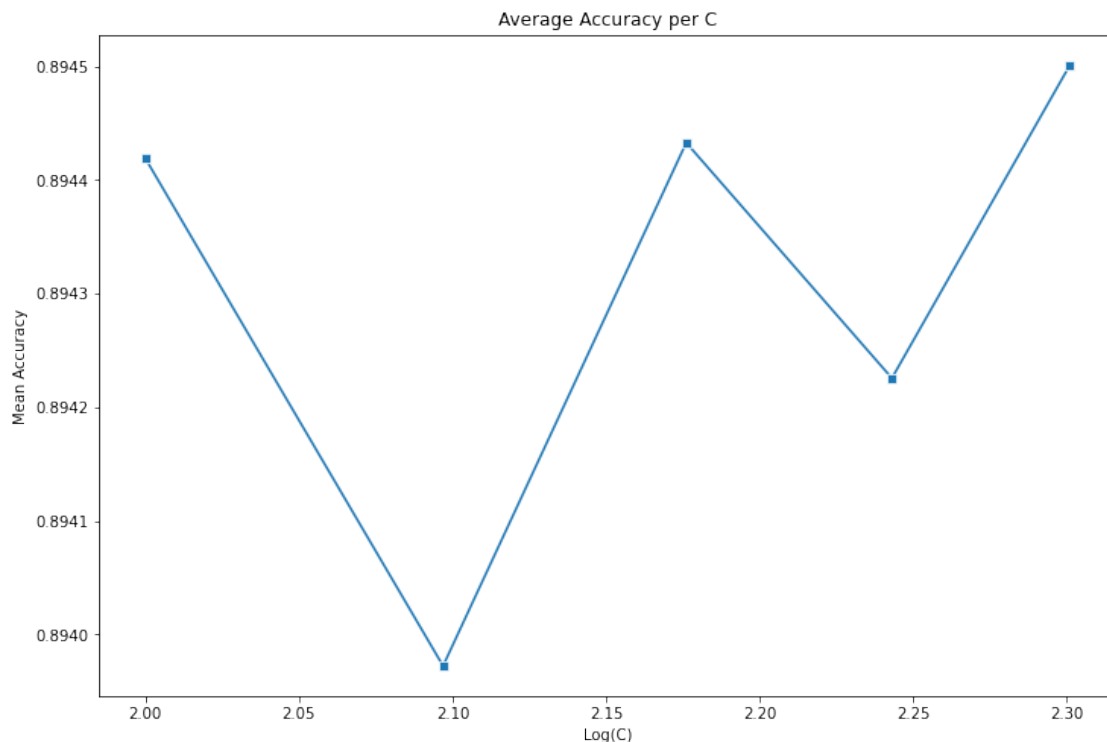
Now that our accuracy is a good bit higher than where we started, we can begin tuning the hyperparameters to see if we can get a better accuracy

```
[27]: C_list = [100, 125, 150, 175, 200]
cv_scores = []

for c in C_list:
    logreg = LogisticRegression(C = c, solver='liblinear', max_iter=10000)
    cv_loop_results = cross_validate(X = X_train_samp, y = y_train_samp,
    ↪ estimator=logreg, cv=10)
    cv_scores.append(np.mean(np.sqrt(np.abs(cv_loop_results['test_score']))))
    print(f"Done with C = {c}")

fig, ax = plt.subplots(figsize=(12, 8))
sns.lineplot(x = np.log10(C_list), y = cv_scores, marker = 's')
ax.set_xlabel('Log(C)')
ax.set_ylabel('Mean Accuracy')
ax.set_title('Average Accuracy per C');
```

```
Done with C = 100
Done with C = 125
Done with C = 150
Done with C = 175
Done with C = 200
```



```
[28]: print(cv_scores)
index = cv_scores.index(max(cv_scores))
print(f"A C of {C_list[index]} is the best performing c with a score of:␣
↪{cv_scores[index]}")
```

```
[0.8944189106752531, 0.8939721736670734, 0.8944327842262105, 0.8942255197636193,
0.8945013355816519]
```

A C of 200 is the best performing c with a score of: 0.8945013355816519

With the cross validation testing finding that a c of 100 performs the best, we can put this back into our model and see if this reflects.

```
[29]: h_model = LogisticRegression(solver='liblinear', C=200, max_iter=10000)
h_model.fit(X_train_imp, y_train)

y_htrain_pred = h_model.predict(X_train_imp)
y_hpred = h_model.predict(X_test_imp)
print(f"Train Acc: {accuracy_score(y_train, y_htrain_pred)}\nTest Acc:␣
↪{accuracy_score(y_test, y_hpred)}")
print(f"Specificity(Recall) Score: {recall_score(y_test, y_hpred,␣
↪pos_label=0)}")
print(f"Recall Score: {recall_score(y_test, y_hpred)}")
```

Train Acc: 0.790280583613917

Test Acc: 0.7908417508417508


Specificity(Recall) Score: 0.5858498705780846

Recall Score: 0.9220320265046935

```
[30]: list_acc.append(accuracy_score(y_test, y_hpred))
list_spec.append(recall_score(y_test, y_hpred, pos_label=0))
list_rec.append(recall_score(y_test, y_hpred))
```

```
[31]: fig, ax = plt.subplots(figsize=(10,8))
fig.set_facecolor('grey')
cm = confusion_matrix(y_test, y_hpred)
disp = ConfusionMatrixDisplay(cm)
disp.plot(ax=ax)
```

```
[31]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x29cbd1feaf0>
```



We may be approaching the limit of what can be done with a logistic regression model. We got a peak test accuracy of around 79 percent. We will now move on to our second model to see if we can improve our score

4 Building a Random Forest classifier

```
[32]: from sklearn.ensemble import RandomForestClassifier

clf = RandomForestClassifier(random_state=42)
clf.fit(X_train_samp, y_train_samp)

train_tree_pred = clf.predict(X_train_samp)
test_tree_pred = clf.predict(X_test_imp)
print(f"Train Acc: {accuracy_score(y_train_samp, train_tree_pred)}\nTest Acc:␣
      ↳{accuracy_score(y_test, test_tree_pred)}")
print(f"Specificity(Recall) Score: {recall_score(y_test, test_tree_pred,␣
      ↳pos_label=0)}")
```

```
print(f"Recall Score: {recall_score(y_test, test_tree_pred)}")
```

Train Acc: 0.9984738926637841


Test Acc: 0.8545454545454545

Specificity(Recall) Score: 0.7891285591026748

Recall Score: 0.896410822749862

```
[33]: fig, ax = plt.subplots(figsize=(10,8))
      cm = confusion_matrix(y_test, test_tree_pred)
      fig.set_facecolor('grey')
      disp = ConfusionMatrixDisplay(cm)
      disp.plot(ax=ax)
```

```
[33]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x29cbd295b80>
```



We can see that this is already considerably better than the previous logistic regression model. However, there is a fair amount of overfitting so we can now tune some hyperparameters to see if we can improve it.

5 Max Tree Depth (RAW)


```
[34]: # Identify the optimal tree depth for given data
from sklearn.metrics import auc, roc_curve
SEED = 42
depths = np.arange(1, 33)
train_results = []
test_results = []
for depth in depths:
    dt_d = RandomForestClassifier(criterion='entropy', random_state=SEED,
    ↪max_depth=depth)
    dt_d.fit(X_train_samp, y_train_samp)

    y_train_pred = dt_d.predict(X_train_samp)
    fpr, tpr, thresholds = roc_curve(y_train_samp, y_train_pred)
    roc_auc = auc(fpr, tpr)
    train_results.append(roc_auc)

    y_test_pred = dt_d.predict(X_test_imp)
    fpr, tpr, thresholds = roc_curve(y_test, y_test_pred)
    roc_auc = auc(fpr, tpr)
    test_results.append(roc_auc)

fig, ax = plt.subplots(figsize=(12, 8))
plt.plot(depths, train_results, 'b', label='Train AUC')
plt.plot(depths, test_results, 'r', label='Test AUC')
plt.legend()
plt.xlabel('Tree Depth')
plt.ylabel('AUC Score')
```

```
[34]: Text(0, 0.5, 'AUC Score')
```



Best max tree depth is around 12-17

6 Minimum Sample Split (RAW)


```
[35]: # Identify the optimal min-samples-split for given data
min_sample_splits = np.linspace(0.001, 0.01, 10)
train_results = []
test_results = []
for sample in min_sample_splits:
    dt_d = RandomForestClassifier(criterion='entropy', random_state=SEED,
    ↪min_samples_split=sample)
    dt_d.fit(X_train_samp, y_train_samp)

    y_train_pred = dt_d.predict(X_train_samp)
    fpr, tpr, thresholds = roc_curve(y_train_samp, y_train_pred)
    roc_auc = auc(fpr, tpr)
    train_results.append(roc_auc)

    y_test_pred = dt_d.predict(X_test_imp)
    fpr, tpr, thresholds = roc_curve(y_test, y_test_pred)
    roc_auc = auc(fpr, tpr)
    test_results.append(roc_auc)
```

```
fig, ax = plt.subplots(figsize=(12, 8))
plt.plot(min_sample_splits, train_results, 'b', label='Train AUC')
plt.plot(min_sample_splits, test_results, 'r', label='Test AUC')
plt.legend()
plt.xlabel('Min Sample Splits')
plt.ylabel('AUC Score')
```

```
[35]: Text(0, 0.5, 'AUC Score')
```



Best minimum sample split is around 0.1

7 Minimum Sample Leafs (RAW)

```
[36]: # Calculate the optimal value for minimum sample leafs
min_sample_leafs = np.linspace(0.0001, 0.001, 10)
train_results = []
test_results = []
for leafs in min_sample_leafs:
    dt_d = RandomForestClassifier(criterion='entropy', random_state=SEED,
    ↪ min_samples_leaf=leafs)
    dt_d.fit(X_train_samp, y_train_samp)

    y_train_pred = dt_d.predict(X_train_samp)
```

```


fpr, tpr, thresholds = roc_curve(y_train_samp, y_train_pred)
roc_auc = auc(fpr, tpr)
train_results.append(roc_auc)

y_test_pred = dt_d.predict(X_test_imp)
fpr, tpr, thresholds = roc_curve(y_test, y_test_pred)
roc_auc = auc(fpr, tpr)
test_results.append(roc_auc)

fig, ax = plt.subplots(figsize=(12, 8))
plt.plot(min_sample_leafs, train_results, 'b', label='Train AUC')
plt.plot(min_sample_leafs, test_results, 'r', label='Test AUC')
plt.legend()
plt.xlabel('Min Sample Leafs')
plt.ylabel('AUC Score')

```

[36]: Text(0, 0.5, 'AUC Score')



The best minimum sample leaf is around 0.1

8 Maximum Features (RAW)

```
[37]: # Find the best value for optimal maximum feature size
max_features = np.arange(14, 20)
test_results = []
train_results = []


for feature in max_features:
    dt = RandomForestClassifier(criterion='entropy', random_state=SEED,
    ↪max_features=feature)
    dt.fit(X_train_samp, y_train_samp)

    y_train_pred = dt.predict(X_train_samp)
    fpr, tpr, thresholds = roc_curve(y_train_samp, y_train_pred)
    auc_score = auc(fpr, tpr)
    train_results.append(auc_score)

    y_test_pred = dt.predict(X_test_imp)
    fpr, tpr, thresholds = roc_curve(y_test, y_test_pred)
    auc_score = auc(fpr, tpr)
    test_results.append(auc_score)

fig, ax = plt.subplots(figsize=(12, 8))
plt.plot(max_features, test_results, 'r', label='Test AUC')
plt.plot(max_features, train_results, 'b', label='Train AUC')
plt.legend()
plt.ylabel('AUC Score')
plt.xlabel('Max Features')
```

```
[37]: Text(0.5, 0, 'Max Features')
```



The best maximum features count is around 19

With all the features optimized, we took a slight hit to accuracy, but managed to bring the test and train accuracy more inline with each other. However, the specificity score is considerably lower than the recall score. In this case, we deemed it to be more beneficial to get a false negative than a false positive: a broken well being predicted as functional would be worse than a functioning well being predicted as broken. We will try to increase specificity, even at the cost of recall and overall accuracy to a limit.

```
[38]: from sklearn.ensemble import RandomForestClassifier

#clf = RandomForestClassifier(random_state=42, min_samples_leaf=15,
#                             max_depth=20, max_features=19, min_samples_split=0.001, class_weight={0: 1,
#                             1: 0.75})
clf = RandomForestClassifier(random_state=42, min_samples_leaf=20,
                             max_depth=15, max_features=19, min_samples_split=0.001)
#clf = RandomForestClassifier(random_state=42)
clf.fit(X_train_samp, y_train_samp)

train_tree_pred = clf.predict(X_train_samp)
test_tree_pred = clf.predict(X_test_imp)
print(f"Train Acc: {accuracy_score(y_train_samp, train_tree_pred)}\nTest Acc:
      {accuracy_score(y_test, test_tree_pred)}")
```

```
print(f"Specificity(Recall) Score: {recall_score(y_test, test_tree_pred, pos_label=0)}")
print(f"Recall Score: {recall_score(y_test, test_tree_pred)}")
```

Train Acc: 0.8546019403364703

Test Acc: 0.8325925925925926


Specificity(Recall) Score: 0.7111302847282139

Recall Score: 0.9103257868580894

```
[39]: list_acc.append(accuracy_score(y_test, test_tree_pred))
      list_spec.append(recall_score(y_test, test_tree_pred, pos_label=0))
      list_rec.append(recall_score(y_test, test_tree_pred))
```

```
[40]: fig, ax = plt.subplots(figsize=(10,8))
      cm = confusion_matrix(y_test, test_tree_pred)
      fig.set_facecolor('grey')
      disp = ConfusionMatrixDisplay(cm)
      disp.plot(ax=ax)
```

```
[40]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x29cba232400>
```



We can specify class weights to make the model less likely to predict positive.

```
[41]: clf = RandomForestClassifier(random_state=42, min_samples_leaf=20,
    ↪max_depth=15, max_features=19, min_samples_split=0.001, class_weight={0: 1,
    ↪1: 0.65})
clf.fit(X_train_samp, y_train_samp)

train_tree_pred = clf.predict(X_train_samp)
test_tree_pred = clf.predict(X_test_imp)
print(f"Train Acc: {accuracy_score(y_train_samp, train_tree_pred)}\nTest Acc:
    ↪{accuracy_score(y_test, test_tree_pred)}")
print(f"Specificity(Recall) Score: {recall_score(y_test, test_tree_pred,
    ↪pos_label=0)}")
print(f"Recall Score: {recall_score(y_test, test_tree_pred)}")
```

Train Acc: 0.8514043821082082

Test Acc: 0.8112457912457912

Specificity(Recall) Score: 0.8200172562553926


Recall Score: 0.8056322473771397

```
[42]: list_acc.append(accuracy_score(y_test, test_tree_pred))
list_spec.append(recall_score(y_test, test_tree_pred, pos_label=0))
list_rec.append(recall_score(y_test, test_tree_pred))
```

This is around the best we can do with train and test accuracy are more in line with each other. We did take a hit to both training and test acc, but since they are closer together, there should be less variance when tested on other data.

```
[43]: fig, ax = plt.subplots(figsize=(10,8))
cm = confusion_matrix(y_test, test_tree_pred)
fig.set_facecolor('grey')
disp = ConfusionMatrixDisplay(cm)
disp.plot(ax=ax)
```


```
[43]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x29cc26f8e50>
```

9 Line chart of accuracy over time

```
[44]: fig, ax = plt.subplots(figsize=(12, 8))
x_arr = ['Baseline', 'Improved Basline', 'Random Forest', 'Spec Random Forest']
plt.plot(x_arr, list_acc, label='Test Accuracy', c='g', marker='s')
plt.plot(x_arr, list_spec, label='Specificity Score', c='b', marker='s')
plt.plot(x_arr, list_rec, label='Recall Score', c='r', marker='s')
ax.set_facecolor('grey')
fig.set_facecolor('grey')

plt.legend()
plt.grid()
plt.xlabel('Model')
plt.ylabel('Accuracy/Score');
```



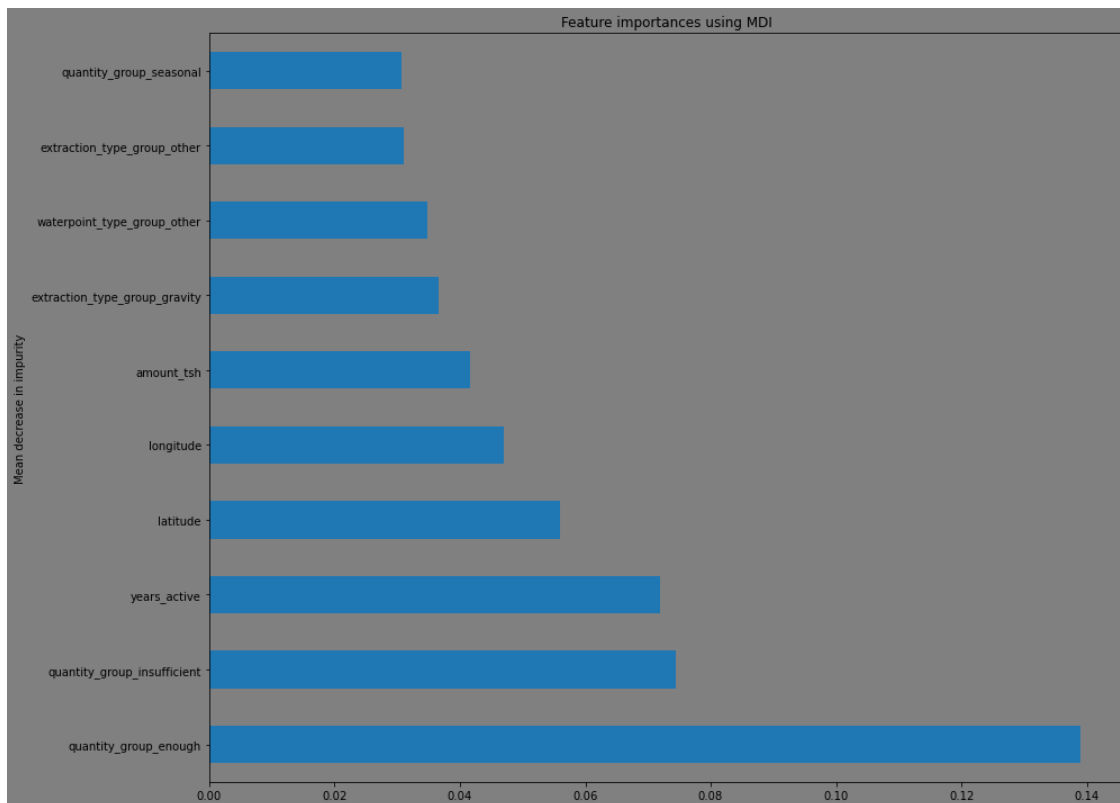
10 Feature Importance

Next we will try to determine which features played a bigger role in determining the target

```
[45]: #Mean decrease in impurity
importances = clf.feature_importances_
names = X_train_samp.columns
mdi = {k:v for (k,v) in zip(names, importances)}

#sort by importance values
#top_20_mdi = sorted(mdi.items(), key=lambda item: item[1], reverse=True)[:20]
mdi_ser = pd.Series(mdi)
top_20_mdi = mdi_ser.sort_values(ascending=False)[:10]
```

```
[46]: fig, ax = plt.subplots(figsize=(14, 10))
top_20_mdi.plot.barh(ax=ax)
fig.set_facecolor('grey')
ax.set_facecolor('grey')
ax.set_title("Feature importances using MDI")
ax.set_ylabel("Mean decrease in impurity")
fig.tight_layout()
#fig.autofmt_xdate()
```



Here when we sort by the importance values, we see some interesting categories. Since we one hot encoded a lot of the categories, we will choose specific attributes of the column that are deemed important.

```
[47]: #Feature permutation
from sklearn.inspection import permutation_importance

result = permutation_importance(clf, X_test_imp, y_test, n_repeats=5,
    ↪random_state=42, n_jobs=2, scoring='accuracy')
forest_importances = pd.Series(result.importances_mean, index=names)
```


```
[48]: top_20_imp = forest_importances.sort_values(ascending=False)[:10]
```

```
[49]: from matplotlib import rcParams

fig, ax = plt.subplots(figsize=(14, 10))
top_20_imp.plot.barh(ax=ax)

rcParams['font.weight'] = 'bold'
fig.set_facecolor('grey')
ax.set_facecolor('grey')
ax.set_title("Feature importances using permutation on full model")
```

```
ax.set_ylabel("Mean decrease in impurity")
#fig.autofmt_xdate()
fig.tight_layout()
```



Here we see that within the top 10 of mean difference in impurity and permutation feature importance, `quantity_group_enough` insufficient and seasonal, `years_active`, `latitude`, and `extraction_type_group_gravity` all appear. Thus, we will deem that these factors are important in deciding whether or not a well is functional.